

Validation-Swap Distributions

Diagnosing Machine-Scored Outcomes in Experiments

August 2026

Abstract

In an experiment with clustered treatment assignment, a machine-learning or large language model scores every outcome, and human labels are affordable for only some clusters. Had assignment been individual, fixed scoring error would be balanced by design and only treatment-dependent error would bias the contrast; clustering removes that protection, leaving few effective units for both balancing and validation. Before coding begins, the researcher must decide how many clusters to validate, and the decision requires knowing what each budget buys. Two-phase and prediction-powered estimators correct the treatment effect with the labeled clusters, and power analysis prices a labeling design by the estimator’s asymptotic variance. When scoring error is concentrated in a few clusters, the variance misprices the risk: the corrected estimate’s distribution across possible label sets is discrete and far from Normal, with most draws moving the estimate not at all and a few moving it far. We price the design by that distribution itself. For each budget k we swap human labels into a set of k clusters, recompute the treatment effect, and enumerate all sets exactly, which is feasible in benchmarks, pilots, and fully coded studies used to plan smaller budgets. As k grows, the mean of the distribution moves on a straight line from the machine-scored estimate to the fully human-coded one, whatever the error structure, so everything that distinguishes labeling designs lies in the dispersion and tails. In a worked example, two scoring systems agree on RMSE, R^2 , both endpoints, and the entire mean path, yet when the error sits in two of twenty clusters, labeling five clusters leaves the estimate unchanged 55.3% of the time. With a single real validation draw the full distribution is out of reach, but its variance at every budget remains estimable without bias, and a negligibility condition says when a Normal approximation can be trusted: only when no small set of clusters dominates the correction. The distribution conditions on the observed pool of gold labels; inference about the population still requires a probability validation design or a transport model.

Researchers can score every response in an experiment with a classifier, a computer-vision model, or a large language model while paying human coders to score only a fraction (Gilardi et al., 2023). The cheap score may have excellent test-set accuracy and still distort the treatment effect. The relevant error is not average scoring error. It is the difference in scoring error between treatment arms.

Suppose an experiment tests whether a message changes the hostility expressed in an open-ended response. A language model scores hostility for all respondents, and human coders score a random subset. The model can overstate hostility by two hundredths in one arm and understate it by two hundredths in the other. Its mean error is zero, but the estimated treatment effect is wrong by four hundredths. Aggregate accuracy conceals this error because accuracy and the treatment contrast weight observations differently.

Existing methods already correct downstream inference with a small set of gold labels. Post-prediction inference, design-based supervised learning, and prediction-powered inference combine predictions on many observations with gold-label residuals on a smaller probability sample (Wang et al., 2020; Egami et al., 2023; Angelopoulos et al., 2023a,b). Two-phase designs for allocating expensive measurements go back to Neyman (1938); modern treatments choose the validation sample to minimize the variance of a target estimator (Amorim et al., 2021), including adaptively (Zrnic and Candès, 2024). Parallel corrections exist when the machine prediction enters the analysis as a regression covariate rather than the outcome (Fong and Tyler, 2021). Work on text as data explains why a learned outcome requires separate measurement and causal design choices (Egami et al., 2022). These methods establish the estimator and its uncertainty.

This paper asks how much the corrected answer depends on which clusters receive the gold labels, and builds the argument from the case where the answer is boring. In an individually randomized experiment, scoring error that is a fixed attribute of the response is balanced by randomization like any baseline covariate, however strongly it clusters by dialect, site, or coder batch. Only error that depends on treatment itself biases the contrast, and an ordinary validation sample corrects that bias with textbook two-phase theory. Clustered assignment removes this protection twice over: the balancing now operates over a few clusters rather than many individuals, and validation samples those same few clusters, so the correction inherits the randomness of which clusters were checked. We characterize that randomness exactly through the fixed-budget random-swap distribution of the treatment-effect estimate. Its spread shows whether a validation budget produces a stable correction or a lottery; its tails show how often the estimate moves by more than a prespecified substantive amount.

The contribution is deliberately narrow. The estimator below is a special case of established difference and prediction-powered estimators, and each formal result is a short application of standard sampling theory. What is new is the object: the exact distribution of the corrected estimate over labeling designs, used to price a validation budget before it is spent.

1 Machine scores as experimental outcomes

An experiment measures an outcome for units $i = 1, \dots, N$, of which N_1 are treated and N_0 are control. Human coding of the outcome Y_i is expensive, so a machine score $\hat{Y}_i$ stands in for it. Define the scoring error

$$e_i = Y_i - \hat{Y}_i. \quad (1)$$

The machine score must be fixed independently of the outcome analysis. A score learned on the experimental outcomes requires a separate training sample or an appropriate cross-fitting design. Swaps do not repair overfitting.

For clarity, consider the unadjusted difference in means. The proxy and gold estimates are

$$\hat{\tau}_P = \frac{1}{N_1} \sum_{i:D_i=1} \hat{Y}_i - \frac{1}{N_0} \sum_{i:D_i=0} \hat{Y}_i, \quad (2)$$

$$\hat{\tau}_G = \frac{1}{N_1} \sum_{i:D_i=1} Y_i - \frac{1}{N_0} \sum_{i:D_i=0} Y_i, \quad (3)$$

and their gap is a weighted sum of the errors,

$$\hat{\tau}_G - \hat{\tau}_P = \sum_{i=1}^N w_i e_i, \quad w_i = \frac{D_i}{N_1} - \frac{1 - D_i}{N_0}. \quad (4)$$

Whether scoring error is a problem is entirely a question about this weighted sum. Error orthogonal to the treatment weights leaves the contrast intact no matter how large the errors are.

Like outcomes, errors have potential values: $e_i(d) = Y_i(d) - \hat{Y}_i(d)$ for assignment $d \in \{0, 1\}$. Call the error *non-differential* when $e_i(1) = e_i(0)$, so the score misreads unit i 's response by the same amount under either assignment. *Differential* error, such as a model that misreads the style treated respondents adopt, violates this condition.

Sections 2 and 3 take expectations over treatment assignment to establish orders of magnitude. From Section 4 onward the analysis conditions on the realized experiment: outcomes, scores, and assignment are fixed, and the only randomness studied is the validation draw.

2 The base case: individually randomized experiments

Assign treatment by complete randomization at the individual level. Under non-differential error, each e_i is a fixed attribute of unit i , and the gap in equation (4) averages fixed numbers over a random assignment:

$$\mathbb{E}_R[\hat{\tau}_G - \hat{\tau}_P] = 0, \quad \mathbb{V}_R[\hat{\tau}_G - \hat{\tau}_P] = \left(\frac{1}{N_1} + \frac{1}{N_0} \right) S_e^2, \quad (5)$$

where S_e^2 is the finite-population variance of the errors and $\mathbb{E}_R$ averages over assignments. Scoring error behaves exactly like a baseline covariate: randomization balances it, and its realized imbalance is $O_p(N^{-1/2})$ noise of the same order as the sampling variability the standard error already reports.

Group structure in the errors changes nothing. If a group h of n_h units — a dialect, a site, a coder batch — shares a common error component μ_h , its contribution to the gap is

$$\mu_h \left(\frac{n_{h1}}{N_1} - \frac{n_{h0}}{N_0} \right), \quad (6)$$

where n_{h1} and n_{h0} count the group's treated and control units. Under individual randomization the imbalance in parentheses is $O_p(\sqrt{n_h}/N)$: group-correlated error cancels across arms because randomization splits every group.

The one threat that survives the base case is differential error. If $e_i(1) \neq e_i(0)$, the gap converges to the average error contrast $\mathbb{E}[e(1) - e(0)]$, a bias no sample size repairs. Even then nothing bespoke is required: a simple random sample of m units receives human labels, the paired discrepancies estimate the bias at rate $m^{-1/2}$, and the standard two-phase and prediction-powered corrections apply with independent-sampling theory (Neyman, 1938; Angelopoulos et al., 2023a).

What cannot substitute for those labels is an accuracy audit.

Proposition 1 (Sharp accuracy bound). *Let $r = (N^{-1} \sum_i e_i^2)^{1/2}$ be the root mean squared error of*

the machine score on the experimental sample. For every error configuration,

$$|\hat{\tau}_G - \hat{\tau}_P| \leq r \sqrt{N \left(\frac{1}{N_1} + \frac{1}{N_0} \right)}, \quad (7)$$

with equality exactly when e_i is proportional to the contrast weight w_i .

Proof. The gap in equation (4) is linear in the errors. By the Cauchy–Schwarz inequality its absolute value is at most $\|w\| \|e\|$, where $\|w\|^2 = 1/N_1 + 1/N_0$ and $\|e\|^2 = Nr^2$, with equality exactly when e is proportional to w . $\square$

The bound is attainable, and it is weak at the scale that matters. In the numerical design below, RMSE 0.200 permits an endpoint gap as large as 0.400, 4 times the gold treatment effect of 0.100. Accuracy metrics constrain the contrast only at the order of r itself, and they cannot separate differential from non-differential error. That separation requires labels.

3 What clustered assignment breaks

Now assign treatment at the cluster level, as in a site-, classroom-, or wave-randomized experiment: clusters $g = 1, \dots, G$ with members indexed by i , treatment D_g constant within cluster, and errors written e_{gi} . Three things break at once.

First, the balancing in equation (5) now operates over G clusters rather than N individuals. Over cluster randomization with equal cluster sizes and equal arms, the gap is still mean zero, but its standard deviation is $2S_\mu/\sqrt{G}$, where S_μ is the between-cluster standard deviation of mean scoring error. With twenty clusters and cluster-level error variation comparable to the effects under study, the realized experiment carries a gap of the same order as the estimand. What was vanishing noise is now a fixed, consequential number.

Second, clusters are maximal-imbalance groups. Each cluster sits wholly in one arm, so the imbalance factor in equation (6) takes its extreme value for every cluster-shaped error group, and nothing cancels within arms. Define cluster g ’s contribution to the endpoint gap as

$$\psi_g = \begin{cases} N_1^{-1} \sum_i e_{gi}, & D_g = 1, \\ -N_0^{-1} \sum_i e_{gi}, & D_g = 0. \end{cases} \quad (8)$$

The contributions add exactly:

$$\sum_{g=1}^G \psi_g = \hat{\tau}_G - \hat{\tau}_P. \quad (9)$$

Equation (9) is an accounting identity within the paired gold-label pool. It does not identify the errors of unlabeled clusters. For OLS with fixed covariates, residualized-treatment weights replace the treated and control weights in equation (8), and the same adding-up and sampling identities follow; a nonlinear estimator can still be recomputed over swap sets, but its mean path need not be linear. When scoring failures are cluster-shaped — a model that misreads one site’s context is wrong for that whole site — the gap concentrates in a few large ψ_g .

Third, validation must draw whole clusters, because human recoding is organized in the same batches as the scoring failures. The correction computed from k labeled clusters is therefore a statistic of k draws from a population of G numbers that a few of them may dominate: few effective units, and no central limit comfort when the contributions concentrate. This is the setting the rest of the paper prices. The same construction applies to documents, waves, or any other sampling unit, provided the swap unit matches the validation design and contains the error dependence.

4 The fixed-budget swap distribution

Let S_k be a simple random sample without replacement of exactly k clusters. The raw-swap estimate replaces machine scores by human outcomes in the selected clusters:

$$\hat{\tau}_R(S_k) = \hat{\tau}_P + \sum_{g \in S_k} \psi_g. \quad (10)$$

The distribution of $\hat{\tau}_R(S_k)$ over all $\binom{G}{k}$ sets is the raw fixed-budget swap distribution. It describes partial repair of the observed data. A fixed k is preferable to independent Bernoulli swaps when k represents a coding budget.

Let $\bar{\psi} = G^{-1} \sum_g \psi_g$ and $S_\psi^2 = (G-1)^{-1} \sum_g (\psi_g - \bar{\psi})^2$. Standard sampling identities give

$$\mathbb{E}_{S_k}[\hat{\tau}_R(S_k)] = \hat{\tau}_P + \frac{k}{G}(\hat{\tau}_G - \hat{\tau}_P), \quad (11)$$

$$\mathbb{V}_{S_k}[\hat{\tau}_R(S_k)] = k \left(1 - \frac{k}{G}\right) S_\psi^2. \quad (12)$$

Equation (11) is the central result. In a linear outcome analysis, the average raw-swap path must be a straight line connecting the proxy and gold estimates. Any two scoring systems with the same endpoints have the same mean path. The path says nothing about whether the correction is spread evenly across clusters or concentrated in a few. Equation (12) shows where that information resides. Two further facts follow directly from equations (9) and (11) and need no proof beyond arithmetic. The $k=1$ distribution is uniform on $\{\hat{\tau}_P + \psi_g\}$, so it identifies the multiset of cluster contributions, and every fixed- k distribution is a symmetric function of that multiset. The mean path, by contrast, depends on the contributions only through their sum: any two error patterns with the same endpoints share the entire mean path, however differently the correction is distributed across clusters.

The distribution supports several descriptive summaries. A researcher can report its central quantiles, the chance that swapping k clusters changes the estimate's sign, and the chance that the movement exceeds a prespecified substantive threshold. These probabilities refer to the stated swap design. They are not p -values.

Raw replacement mechanically shrinks as k/G falls. If the goal is to estimate the fully human-coded contrast from a k -cluster probability sample, scale the observed correction by its inclusion probability:

$$\hat{\tau}_H(S_k) = \hat{\tau}_P + \frac{G}{k} \sum_{g \in S_k} \psi_g. \quad (13)$$

This is a Horvitz-Thompson difference correction (Horvitz and Thompson, 1952; Särndal et al., 1992). In this simple setting it is also the rectifier used by prediction-powered inference (Angelopoulos et al., 2023a). It satisfies

$$\mathbb{E}_{S_k}[\hat{\tau}_H(S_k)] = \hat{\tau}_G, \quad (14)$$

$$\mathbb{V}_{S_k}[\hat{\tau}_H(S_k)] = \left(\frac{G}{k}\right)^2 k \left(1 - \frac{k}{G}\right) S_\psi^2. \quad (15)$$

Thus the design-scaled swap distribution is centered on the gold endpoint. Its spread measures validation-sample fragility at budget k .

Unequal-probability validation replaces G/k with inverse inclusion probabilities. Stratifying the validation draw by treatment arm or by an observable proxy-risk score can improve precision. Those choices change the swap distribution and must be included in its definition. The optimal choice depends on the target estimator (Amorim et al., 2021; Zrnic and Candès, 2024).

An exact distribution requires ψ_g for every cluster in the finite pool. That is available in a benchmark dataset, a simulation, or a large pilot used to study smaller future budgets. With only one small validation sample, the analyst cannot recover the unseen cluster contributions nonparametrically. Standard design-based variance estimation remains available when selection probabilities are known. A cluster bootstrap or a model for e_{gi} can approximate more of the distribution, but then the result depends on resampling or transport assumptions.

The second moment of the swap distribution is nevertheless estimable from a single draw, at the observed budget and at every other one.

Proposition 2 (The variance path is estimable from one validation sample). *Let S_k be a simple random sample of $k \geq 2$ clusters, and let $s_\psi^2 = (k-1)^{-1} \sum_{g \in S_k} (\psi_g - \bar{\psi}_{S_k})^2$ be the sample variance of the observed contributions. Then $\mathbb{E}_{S_k}[s_\psi^2] = S_\psi^2$. Substituting s_ψ^2 into equations (12) and (15) therefore gives unbiased estimates of the raw and design-scaled swap variances at every budget $k' \in \{1, \dots, G\}$, and Chebyshev’s inequality bounds the tail probability $\Pr\{|\hat{\tau}_H(S_{k'}) - \hat{\tau}_G| > \delta\}$ by $(G/k')^2 k' (1 - k'/G) S_\psi^2 / \delta^2$ for every threshold $\delta > 0$.*

Proof. Design-unbiasedness of the sample variance under simple random sampling without replacement is classical (Cochran, 1977; Li and Ding, 2017). The variance claims substitute this estimate into the exact formulas, and the tail bound applies Chebyshev’s inequality to the design-scaled estimator, which is centered on $\hat{\tau}_G$ by equation (14). $\square$

A pilot at one budget therefore prices fragility at every other budget. This is the finite-population analogue of power analysis for prediction-powered inference, which chooses the labeled sample size from an asymptotic variance (Chen et al., 2026). The estimate is honest but noisy: the sampling variance of s_ψ^2 depends on fourth moments of the contributions, so with few validated clusters and concentrated error the estimated variance path is itself fragile.

Moment summaries also carry an implicit distributional assumption. The following negligibility condition, due to Hájek (1960) and stated here in the form of Theorem 1 of Li and Ding (2017), says when the swap distribution can be treated as Normal.

Proposition 3 (Normal approximations require diffuse contributions). *Embed the design in a sequence of finite populations with $G \rightarrow \infty$ clusters and budgets $k = k(G)$. If*

$$\frac{1}{\min(k, G - k)} \cdot \frac{\max_g (\psi_g - \bar{\psi})^2}{S_\psi^2} \rightarrow 0, \quad (16)$$

then the standardized raw-swap estimate $\{\hat{\tau}_R(S_k) - \mathbb{E}[\hat{\tau}_R(S_k)]\} / \mathbb{V}[\hat{\tau}_R(S_k)]^{1/2}$ converges in distribution to a standard Normal, and the design-scaled estimate, an affine transformation at fixed k , converges with it. The condition fails when a bounded number of clusters carries a nonvanishing share of the centered squared contributions.

Proof. The raw-swap estimate is $\hat{\tau}_P$ plus k times the sample mean of a simple random sample of the finite population $\{\psi_1, \dots, \psi_G\}$, so the claim is the finite population central limit theorem for simple random sampling (Erdős and Rényi, 1959; Hájek, 1960), in the sufficient form of Theorem 1 of Li and Ding (2017). Sharp necessary and sufficient conditions of Lindeberg type are in Hájek (1960). $\square$

The finite-sample diagnostic implied by condition (16) is the largest single-cluster share of the centered squared contributions. In the numerical design below, one cluster carries 45% of that total in the concentrated case and 25% in the offsetting case, and the diffuse case is degenerate. Budget planning that combines Proposition 2 with a Normal approximation is therefore reliable exactly when the correction is diffuse, and misleading in the concentrated case the diagnostic exists to detect. That is the formal reason to examine the exact distribution rather than its first two moments.

5 A worked example: same accuracy, different prices

The numerical example has 20 equal-size clusters with 50 observations each. Half the clusters are treated. The fully human-coded treatment effect is 0.100, and the machine-scored effect is 0.060 in the diffuse and concentrated cases. All three scoring systems have RMSE 0.200 and the same R^2 because they score the same gold outcomes and have the same squared-error total.

The error patterns differ only in their cluster contributions. In the diffuse case, each cluster contributes 0.040/20 to the endpoint gap. In the concentrated case, two clusters contribute 0.040/2 each and the others contribute zero. The offsetting case has two positive and two negative contributions of equal size, so the proxy and gold endpoints agree even though individual swaps can move the estimate.

Table 1: Matched prediction accuracy with different contribution patterns

Error pattern	Gold ATE	Proxy ATE	Gap	RMSE	R^2	Top-two share
Diffuse	0.100	0.060	0.040	0.200	0.965	10%
Concentrated	0.100	0.060	0.040	0.200	0.965	100%
Offsetting	0.100	0.100	0.000	0.200	0.965	50%

Note: Gold ATE and proxy ATE are differences in means. Gap is gold minus proxy. Top-two share is the two largest absolute cluster contributions divided by the sum of absolute contributions. Each row uses the same gold outcomes. Quantities are deterministic for the generated finite population.

The first two patterns share the entire mean swap path, so no average over validation draws can separate them. Every difference between them is distributional.

Table 2 prices each validation budget two ways. Tail risk is the exact share of validation sets whose scaled estimate lands more than 0.050 from the gold endpoint. The Normal price is the same probability computed from a Normal approximation with the correct finite-population variance of equation (15), the summary a variance-based power analysis would use. The two prices agree only in the diffuse case, where every draw picks up the same correction and both are zero at every budget.

Concentrated error breaks the approximation in both directions. At 5 labeled clusters the Normal price is 29% against an exact 5.3%, because 55.3% of the draws contain neither high-error cluster and leave the estimate at the machine-scored value. The exact price is also not monotone in the budget: the distribution’s few atoms cross the threshold at different budgets as the G/k scaling shrinks, so offsetting risk climbs to 62% at 7 labeled clusters before falling, while the Normal price declines smoothly and understates the risk where it is largest. A larger budget can carry more tail risk than a smaller one, and the variance-based price never shows it. Decision reversals price the same way: 10% of single-cluster draws in the offsetting case reverse the sign of the estimated effect, the analogue of the sign-change criterion in dropping-data robustness (Giordano et al., 2026). Bias is omitted from the table because it is zero up to numerical precision for every cell. A narrow distribution does not prove that human labels measure the intended construct. It says only that the specified validation draw estimates the observed gold-label endpoint consistently.

Table 2: Exact design-scaled distributions at selected budgets

Error pattern	Gold clusters	Mean	SD	5th	95th	Tail risk	Normal price
Diffuse	2	0.100	0.000	0.100	0.100	0%	0%
	5	0.100	0.000	0.100	0.100	0%	0%
	10	0.100	0.000	0.100	0.100	0%	0%
	20	0.100	0.000	0.100	0.100	0%	0%
Concentrated	2	0.100	0.083	0.060	0.260	19%	54%
	5	0.100	0.048	0.060	0.220	5%	29%
	10	0.100	0.028	0.060	0.140	0%	7%
	20	0.100	0.000	0.100	0.100	0%	0%
Offsetting	2	0.100	0.123	-0.100	0.300	35%	68%
	5	0.100	0.071	0.020	0.180	57%	48%
	10	0.100	0.041	0.020	0.180	14%	22%
	20	0.100	0.000	0.100	0.100	0%	0%

Note: Mean, SD, and quantiles describe equation (13) over all fixed-size cluster samples. Tail risk is $\Pr(|\hat{\tau}_H - \hat{\tau}_G| > 0.050)$, computed by exact enumeration, not Monte Carlo. Normal price is the same probability under a Normal approximation with the exact finite-population variance of equation (15). Gold ATE equals 0.100 in every scenario.

6 Related subset diagnostics

A thought experiment locates the swap distribution among its neighbors. Suppose every observation had a gold label. The measurement layer would vanish, and the only randomness left would be treatment assignment: over re-randomization, each cluster’s contribution to the estimate enters with a random sign as the cluster changes arm. The wild cluster bootstrap approximates exactly

this sign-flip process, applying independent mean-zero multipliers to cluster score contributions (Cameron et al., 2008), and the first break in Section 3 is that layer’s variance statement for the scoring errors. Now remove the labels for all but k clusters. A second layer of randomness sits on the same cluster contributions: for the linear contrast,

$$\hat{\tau}_R(S_k) - \mathbb{E}_{S_k}[\hat{\tau}_R(S_k)] = \sum_{g=1}^G \left(Z_g - \frac{k}{G} \right) \psi_g,$$

where $Z_g = \mathbf{1}\{g \in S_k\}$ marks the labeled clusters. The two procedures are one construction — recompute the estimate under random reweighting of cluster contributions — instantiated with different multiplier laws for different design layers: signs for assignment, inclusion indicators for labeling. That is why neither substitutes for the other. Swap quantiles are not confidence intervals, because valid uncertainty for the experimental effect must average over the sign layer, and valid uncertainty for the correction must average over the indicator layer actually used to buy labels.

Dropping-data robustness metrics are the third member of the family: they apply adversarially chosen zero weights to individual observations and locate the worst set with a first-order influence approximation, exact enumeration being infeasible (Giordano et al., 2026). The swap construction differs on every choice — the perturbed unit is the gold-label set rather than the data, the weights are drawn by a design rather than chosen by an adversary, and no approximation is needed because the linear contrast makes every swap effect exactly additive in the cluster contributions. The adversarial variant is nevertheless free here: the most adverse k -cluster validation draw is the sum of the k most extreme contributions, read off by sorting the ψ_g .

7 Use and limits

A validation-swap analysis begins with the same choices required by any two-phase design. The researcher states the causal estimand, fixes the machine scoring rule, identifies the unit sampled for human coding, and records every cluster’s inclusion probability. If treatment is assigned by cluster, swapping individual labels while treating them as independent understates the relevant variation.

The primary analysis should report a valid correction estimator and interval, using prediction-powered, design-based supervised learning, or another method whose assumptions fit the design. The swap distribution is an additional diagnostic. Before seeing it, the researcher chooses validation budgets and a substantively meaningful movement threshold. The exhibit then reports quantiles, tail risk, and any decision reversal probability. Accuracy metrics remain useful descriptions of the scoring model, but they do not substitute for this coefficient-specific analysis.

The diagnostic has four limits. First, it cannot reveal the individual errors of unlabeled observations. Second, a distribution computed inside a convenience validation sample need not transport outside it. Third, human labels can contain their own error or define the wrong construct. Fourth, a small number of validated clusters may leave the distribution too poorly estimated to support fine subgroup claims. The proper conclusion in that case is that the validation budget does not resolve the scoring risk.

Machine-scored experiments create two sources of uncertainty: treatment assignment and gold-label sampling. Randomization-based or cluster-robust inference addresses the first. Established two-phase estimators address the second when they correct the machine-scored contrast. The random-

swap distribution makes validation-sample dependence visible by showing how much the conclusion depends on which clusters were checked.

References

- Gustavo Amorim, Ran Tao, Sarah Lotspeich, Pamela A. Shaw, Thomas Lumley, and Bryan E. Shepherd. Two-phase sampling designs for data validation in settings with covariate measurement error and continuous outcome. *Journal of the Royal Statistical Society: Series A*, 184(4):1368–1389, 2021. doi: 10.1111/rssa.12689.
- Anastasios N. Angelopoulos, Stephen Bates, Clara Fannjiang, Michael I. Jordan, and Tijana Zrnic. Prediction-powered inference. *Science*, 382(6671):669–674, 2023a. doi: 10.1126/science.adi6000.
- Anastasios N. Angelopoulos, John C. Duchi, and Tijana Zrnic. PPI++: Efficient prediction-powered inference, 2023b. arXiv:2311.01453.
- A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Bootstrap-based improvements for inference with clustered errors. *Review of Economics and Statistics*, 90(3):414–427, 2008. doi: 10.1162/rest.90.3.414.
- Yiqun T. Chen, Moran Guo, and Shengyi Li. Power analysis for prediction-powered inference, 2026. arXiv:2603.16041.
- William G. Cochran. *Sampling Techniques*. Wiley, New York, 3rd edition, 1977.
- Naoki Egami, Christian J. Fong, Justin Grimmer, Margaret E. Roberts, and Brandon M. Stewart. How to make causal inferences using texts. *Science Advances*, 8(42):eabg2652, 2022. doi: 10.1126/sciadv.abg2652.
- Naoki Egami, Musashi Hinck, Brandon M. Stewart, and Hanying Wei. Using imperfect surrogates for downstream inference: Design-based supervised learning for social science applications of large language models. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- Paul Erdős and Alfréd Rényi. On the central limit theorem for samples from a finite population. *Publications of the Mathematical Institute of the Hungarian Academy of Sciences*, 4:49–61, 1959.
- Christian Fong and Matthew Tyler. Machine learning predictions as regression covariates. *Political Analysis*, 29(4):467–484, 2021. doi: 10.1017/pan.2020.38.
- Fabrizio Gilardi, Meysam Alizadeh, and Maël Kubli. ChatGPT outperforms crowd workers for text-annotation tasks. *Proceedings of the National Academy of Sciences*, 120(30):e2305016120, 2023. doi: 10.1073/pnas.2305016120.
- Ryan Giordano, Rachael Meager, and Tamara Broderick. An automatic finite-sample robustness metric: When can dropping a little data change conclusions? *Philosophical Transactions of the Royal Society A*, 384(2321):20240614, 2026. doi: 10.1098/rsta.2024.0614.
- Jaroslav Hájek. Limiting distributions in simple random sampling from a finite population. *Publications of the Mathematical Institute of the Hungarian Academy of Sciences*, 5:361–374, 1960.

- Daniel G. Horvitz and Donovan J. Thompson. A generalization of sampling without replacement from a finite universe. *Journal of the American Statistical Association*, 47(260):663–685, 1952. doi: 10.1080/01621459.1952.10483446.
- Xinran Li and Peng Ding. General forms of finite population central limit theorems with applications to causal inference. *Journal of the American Statistical Association*, 112(520):1759–1769, 2017. doi: 10.1080/01621459.2017.1295865.
- Jerzy Neyman. Contribution to the theory of sampling human populations. *Journal of the American Statistical Association*, 33(201):101–116, 1938. doi: 10.1080/01621459.1938.10503378.
- Carl-Erik Särndal, Bengt Swensson, and Jan Wretman. *Model Assisted Survey Sampling*. Springer, New York, 1992.
- Siruo Wang, Tyler H. McCormick, and Jeffrey T. Leek. Methods for correcting inference based on outcomes predicted by machine learning. *Proceedings of the National Academy of Sciences*, 117(48):30266–30275, 2020. doi: 10.1073/pnas.2001238117.
- Tijana Zrnic and Emmanuel J. Candès. Active statistical inference, 2024. arXiv:2403.03208.